Implementation: Value Iteration

In the previous concept, you learned about **value iteration**. In this algorithm, each sweep over the state space effectively performs both policy evaluation and policy improvement. Value iteration is guaranteed to find the optimal policy π_* for any finite MDP.

The pseudocode can be found below.

Value Iteration

Input: MDP, small positive number θ
Output: policy $\pi \approx \pi_*$
Initialize V arbitrarily (e.g., $V(s) = 0$ for all $s \in \mathcal{S}^+$)
repeat
    $\Delta \leftarrow 0$
    for $s \in \mathcal{S}$ **do**
        $v \leftarrow V(s)$
        $V(s) \leftarrow \max_{a \in \mathcal{A}(s)} \sum_{s' \in \mathcal{S}, r \in \mathcal{R}} p(s', r | s, a)(r + \gamma V(s'))$
        $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
    end
until $\Delta < \theta$;
$\pi \leftarrow$ **Policy_Improvement**(MDP, V)
return π

Note that the stopping criterion is satisfied when the difference between successive value function estimates is sufficiently small. In particular, the loop terminates if the difference is less than θ for each state. And, the closer we want the final value function estimate to be to the optimal value function, the smaller we need to set the value of θ.

Feel free to play around with the value of θ in your implementation; note that in the case of the FrozenLake environment, values around `1e-8` seem to work reasonably well.

Implementation

encouraged to pay particular attention to Theorem 3.2. Their main result of interest can be summarized as follows:

> Let V^{final} denote the final value function estimate that is calculated by the algorithm. Then it can be shown that V^{final} differs from the optimal value function v_* by at most $\frac{2\theta\gamma}{1-\gamma}$. In other words, for each $s \in \mathcal{S}$,

> $$\max_{s \in \mathcal{S}} \left| V^{\mathrm{final}}(s) - v_*(s) \right| < \frac{2\theta\gamma}{1-\gamma}.$$

Please use the next concept to complete **Part 6: Value Iteration** of `Dynamic_Programming.ipynb`. Remember to save your work!

If you'd like to reference the pseudocode while working on the notebook, you are encouraged to open **this sheet** in a new window.

Feel free to check your solution by looking at the corresponding section in

NEXT