



DDPG: Agent

We are now ready to put together the actor and policy models to build our DDPG agent. Note that we will need two copies of each model - one local and one target. This is an extension of the "Fixed Q Targets" technique from Deep Q-Learning, and is used to decouple the parameters being updated from the ones that are producing target values.

Here is an outline of the agent class:



DDPG: Agent

```

def __init__(self, task):
    self.task = task
    self.state_size = task.state_size
    self.action_size = task.action_size
    self.action_low = task.action_low
    self.action_high = task.action_high

    # Actor (Policy) Model
    self.actor_local = Actor(self.state_size, self.action_size, self.action_low, self.action_high)
    self.actor_target = Actor(self.state_size, self.action_size, self.action_low, self.action_high)

    # Critic (Value) Model
    self.critic_local = Critic(self.state_size, self.action_size)
    self.critic_target = Critic(self.state_size, self.action_size)

    # Initialize target model parameters with local model parameters
    self.critic_target.model.set_weights(self.critic_local.model.get_weights())
    self.actor_target.model.set_weights(self.actor_local.model.get_weights())

    # Noise process
    self.exploration_mu = 0
    self.exploration_theta = 0.15
    self.exploration_sigma = 0.2
    self.noise = OUNoise(self.action_size, self.exploration_mu, self.exploration_theta, self.exploration_sigma)

    # Replay memory
    self.buffer_size = 100000
    self.batch_size = 64
    self.memory = ReplayBuffer(self.buffer_size, self.batch_size)

    # Algorithm parameters
    self.gamma = 0.99 # discount factor
    self.tau = 0.01 # for soft update of target parameters

def reset_episode(self):
    self.noise.reset()
    state = self.task.reset()
    self.last_state = state
    return state

def step(self, action, reward, next_state, done):
    # Save experience / reward
    self.memory.add(self.last_state, action, reward, next_state, done)

    # Learn, if enough samples are available in memory
    if len(self.memory) > self.batch_size:
        experiences = self.memory.sample()
        self.learn(experiences)

    # Roll over last state and action

```



DDPG: Agent

```

"""Returns actions for given state(s) as per current policy."""
state = np.reshape(state, [-1, self.state_size])
action = self.actor_local.model.predict(state)[0]
return list(action + self.noise.sample()) # add some noise for explo
ration

def learn(self, experiences):
    """Update policy and value parameters using given batch of experience
    tuples."""
    # Convert experience tuples to separate arrays for each element (stat
    es, actions, rewards, etc.)
    states = np.vstack([e.state for e in experiences if e is not None])
    actions = np.array([e.action for e in experiences if e is not None]).
    astype(np.float32).reshape(-1, self.action_size)
    rewards = np.array([e.reward for e in experiences if e is not None]).
    astype(np.float32).reshape(-1, 1)
    dones = np.array([e.done for e in experiences if e is not None]).ast
    pe(np.uint8).reshape(-1, 1)
    next_states = np.vstack([e.next_state for e in experiences if e is no
    t None])

    # Get predicted next-state actions and Q values from target models
    # Q_targets_next = critic_target(next_state, actor_target(next_st
    ate))
    actions_next = self.actor_target.model.predict_on_batch(next_states)
    Q_targets_next = self.critic_target.model.predict_on_batch([next_stat
    es, actions_next])

    # Compute Q targets for current states and train critic model (local)
    Q_targets = rewards + self.gamma * Q_targets_next * (1 - dones)
    self.critic_local.model.train_on_batch(x=[states, actions], y=Q_targe
    ts)

    # Train actor model (local)
    action_gradients = np.reshape(self.critic_local.get_action_gradients
    ([states, actions, 0]), (-1, self.action_size))
    self.actor_local.train_fn([states, action_gradients, 1]) # custom tr
    aining function

    # Soft-update target models
    self.soft_update(self.critic_local.model, self.critic_target.model)
    self.soft_update(self.actor_local.model, self.actor_target.model)

def soft_update(self, local_model, target_model):
    """Soft update model parameters."""
    local_weights = np.array(local_model.get_weights())
    target_weights = np.array(target_model.get_weights())

    assert len(local_weights) == len(target_weights), "Local and target m
    odel parameters must have the same size"

    new_weights = self.tau * local_weights + (1 - self.tau) * target_weig
    hts
    target_model.set_weights(new_weights)

```



DDPG: Agent

batches can introduce a lot of variance into the process, so it's better to perform a soft update, controlled by the parameter τ .

[NEXT](#)